



University | School of
of Glasgow | Computing Science

THE AWARDS
2020

UNIVERSITY
OF THE YEAR

Recurrent Neural Networks

Dr. Fani Deligianni,

fani.deligianni@glasgow.ac.uk

Lecturer (Assistant Professor)

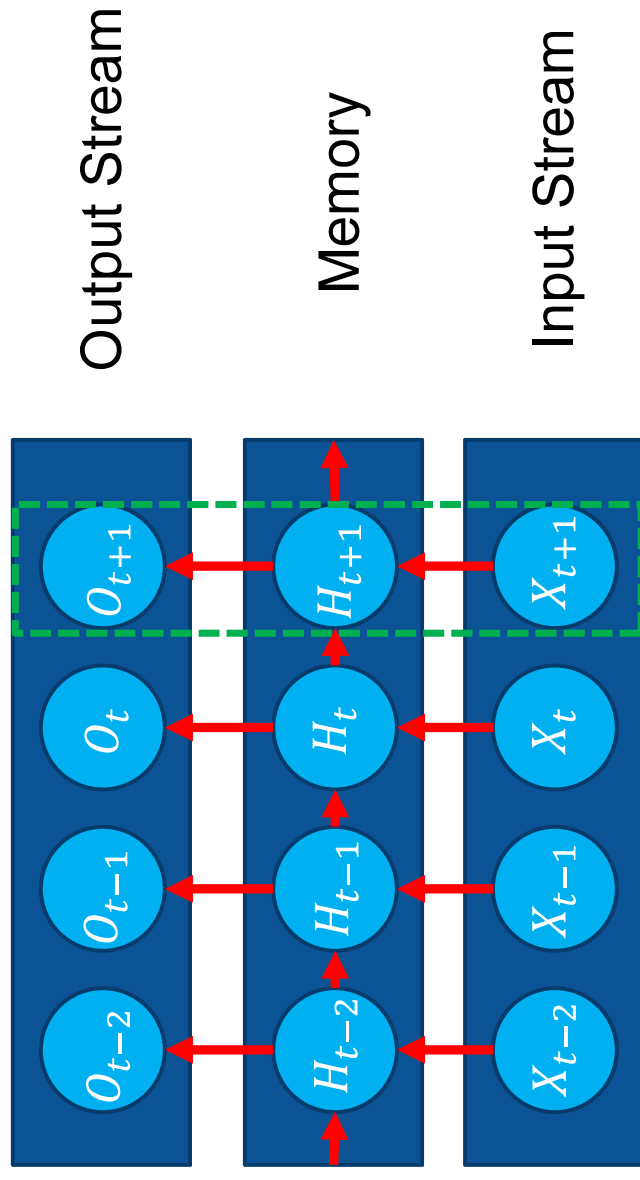
Lead of the Computing Technologies for Healthcare Theme

<https://www.gla.ac.uk/schools/computing/staff/fanideligianni>

WORLD
CHANGING
GLASGOW



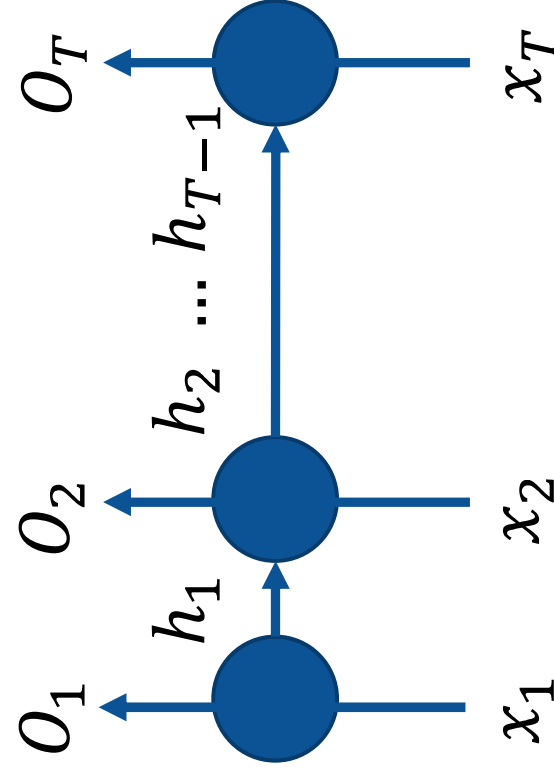
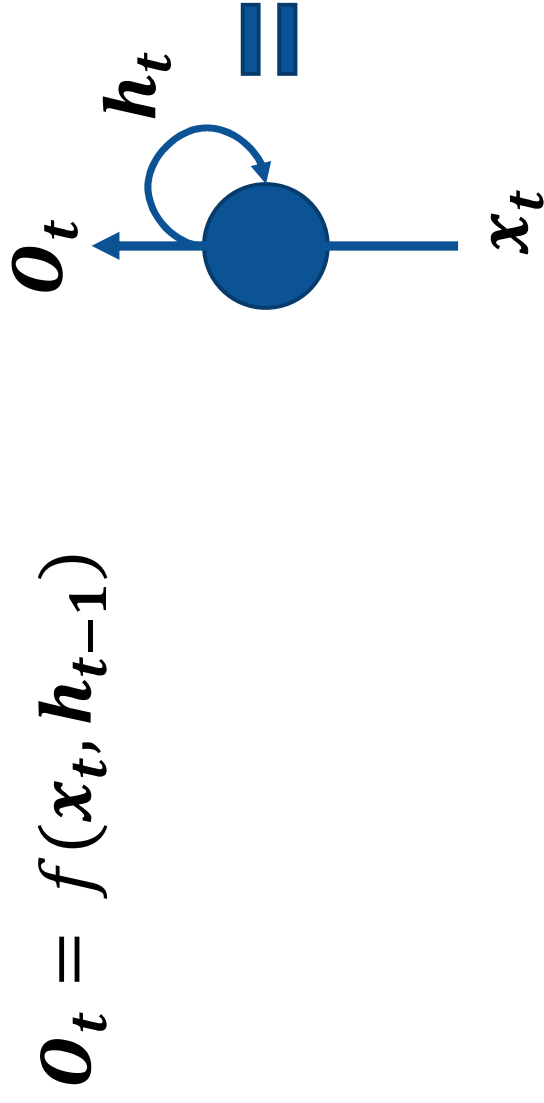
Recurrent Neural Networks



- RNNs accumulate information from each time step
- RNNs can learn representations for variable length sequences



Recurrent Neural Networks



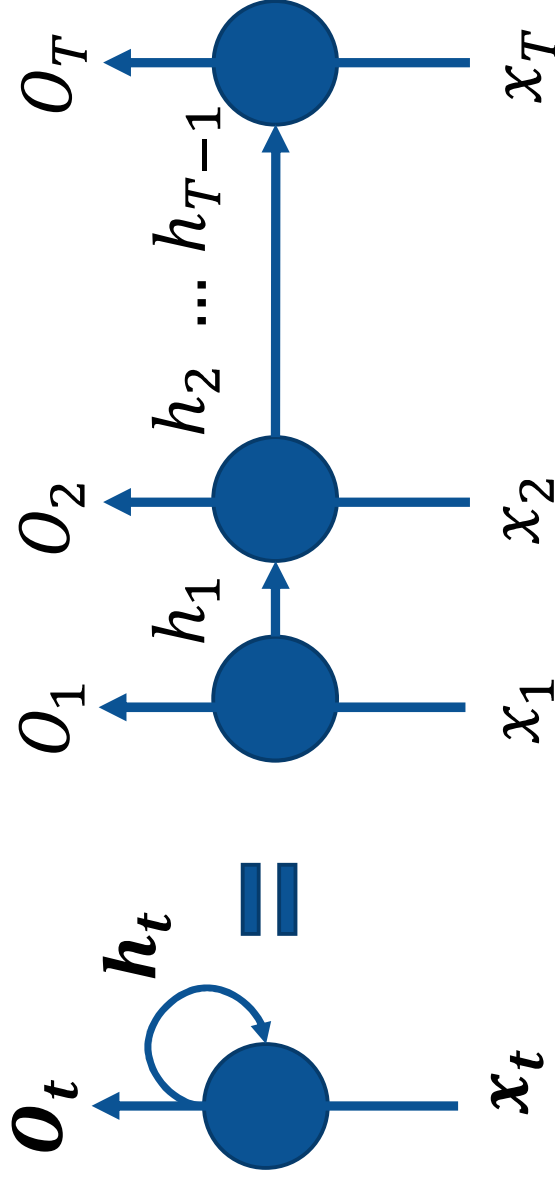
- RNNs accumulate information from each time step
- RNNs can learn representations for variable length sequences



Recurrent Neural Networks

$$O_t = f(x_t, h_{t-1})$$

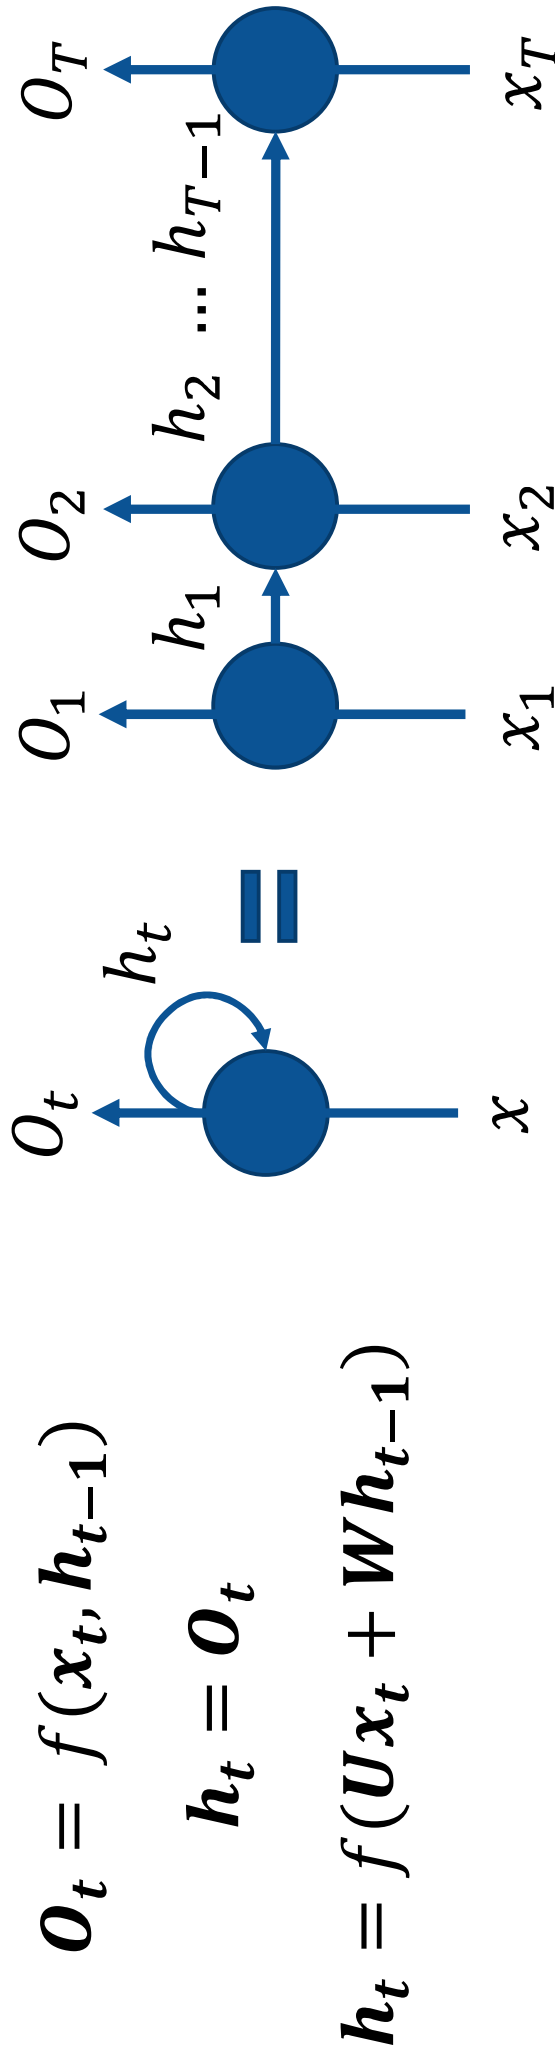
$$h_t = O_t$$



- RNNs accumulate information from each time step
- RNNs can learn representations for variable length sequences



Recurrent Neural Networks



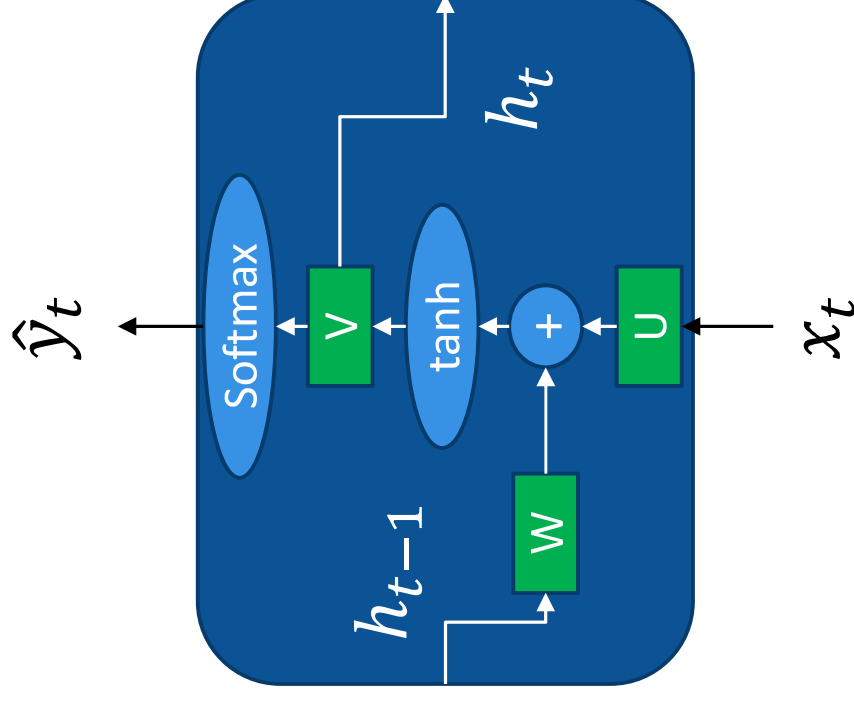
- RNNs accumulate information from each time step
- RNNs can learn representations for variable length sequences



Forward Propagation in RNNs

$$\mathbf{h}_t = \tanh(\mathbf{U}\mathbf{x}_t + \mathbf{W}\mathbf{h}_{t-1})$$

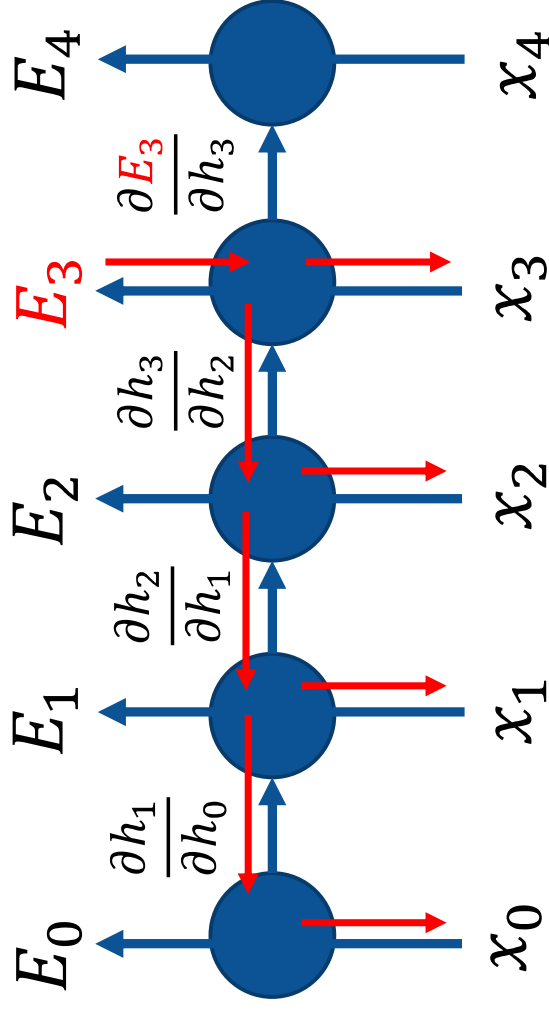
$$\hat{\mathbf{y}}_t = \text{softmax}(\mathbf{V}\mathbf{h}_t)$$



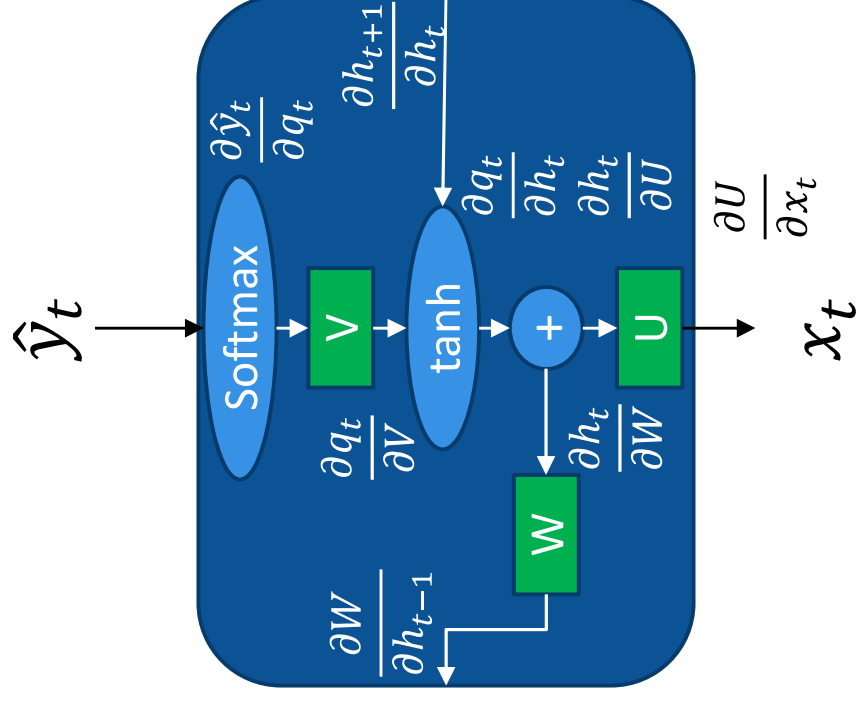
Back Propagation in RNNs

$$E_t = -y_t \log \hat{y}_t$$

$$L(\mathbf{y}, \hat{\mathbf{y}}) = -\frac{1}{N} \sum_t y_t \log \hat{y}_t$$



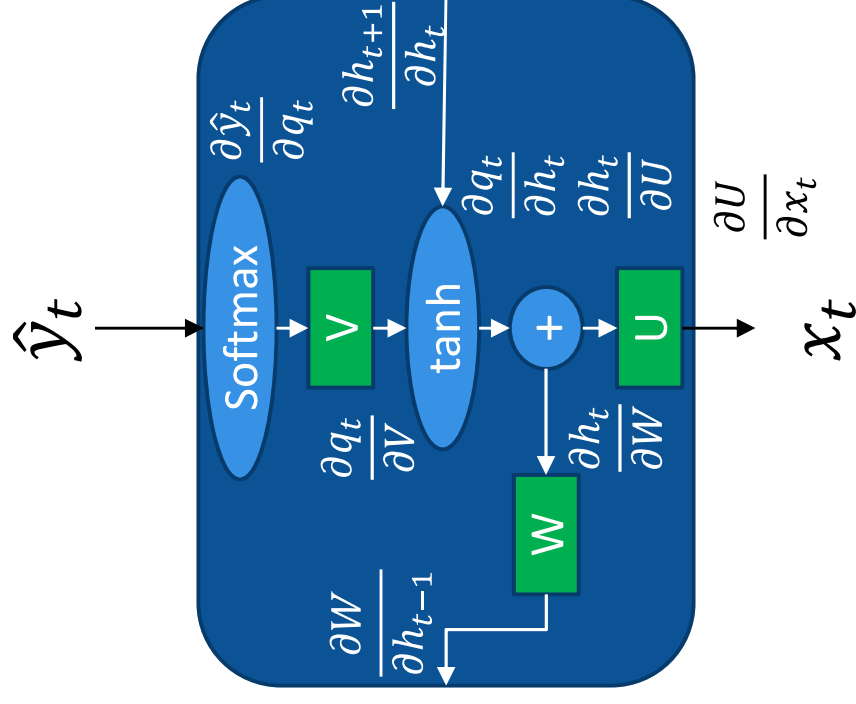
Back Propagation in RNNs



Back Propagation in RNNs

Recursive weights:

$$\frac{\partial E_t}{\partial \mathbf{W}} = \frac{\partial E_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial q_t} \frac{\partial q_t}{\partial h_t} \frac{\partial h_t}{\partial \mathbf{W}}$$

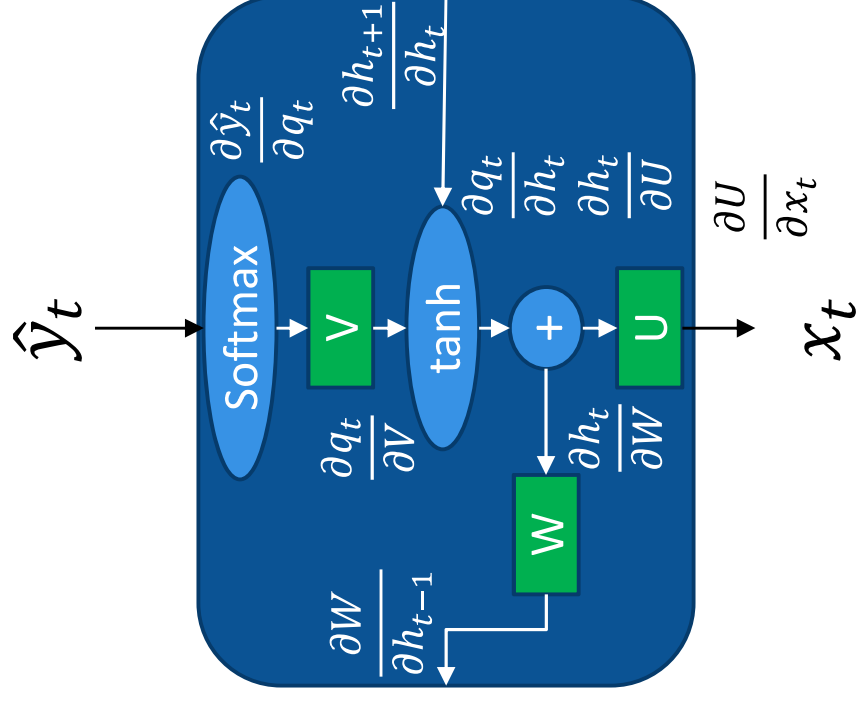


Back Propagation in RNNs

Recursive weights:

$$\frac{\partial E_t}{\partial \mathbf{W}} = \frac{\partial E_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial q_t} \frac{\partial q_t}{\partial h_t} \frac{\partial h_t}{\partial \mathbf{W}}$$

$$\frac{\partial h_t}{\partial \mathbf{W}} \rightarrow \frac{\partial h_t}{\partial \mathbf{W}} + \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial \mathbf{W}}$$



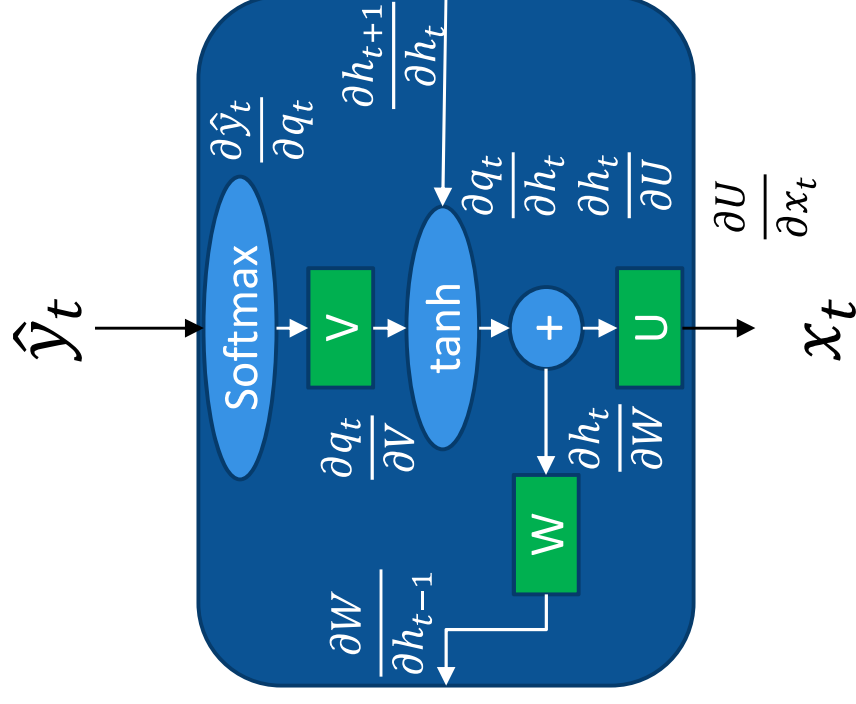
Back Propagation in RNNs

Recursive weights:

$$\frac{\partial E_t}{\partial \mathbf{W}} = \frac{\partial E_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial q_t} \frac{\partial q_t}{\partial h_t} \frac{\partial h_t}{\partial \mathbf{W}}$$

$$\frac{\partial h_t}{\partial \mathbf{W}} \rightarrow \frac{\partial h_t}{\partial \mathbf{W}} + \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial \mathbf{W}}$$

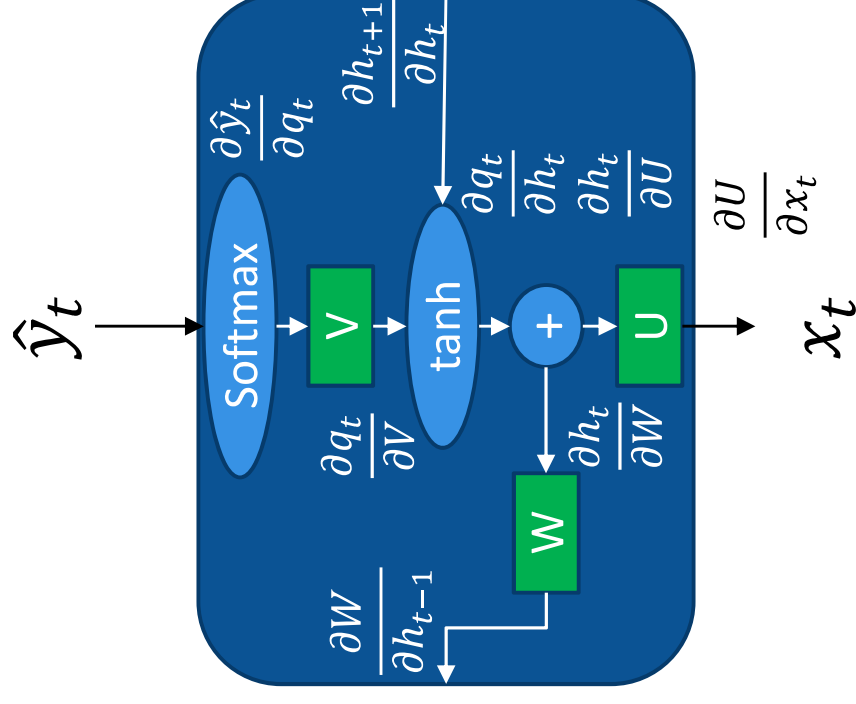
$$\frac{\partial h_{t-2}}{\partial \mathbf{W}}$$



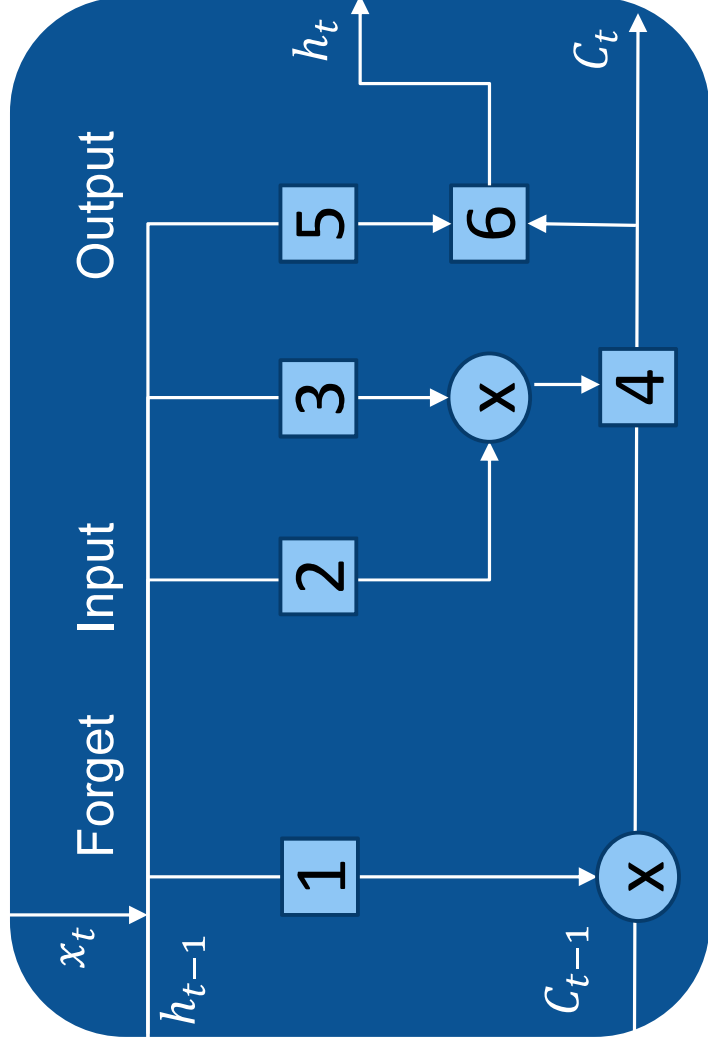
Back Propagation in RNNs

Vanishing/Exploding gradient:

- Careful initialization
- Adaptive rates (Adam optimizer)
- Second order derivatives to predict the occurrence of vanishing gradient



Long Short-Term Memory (LSTM)



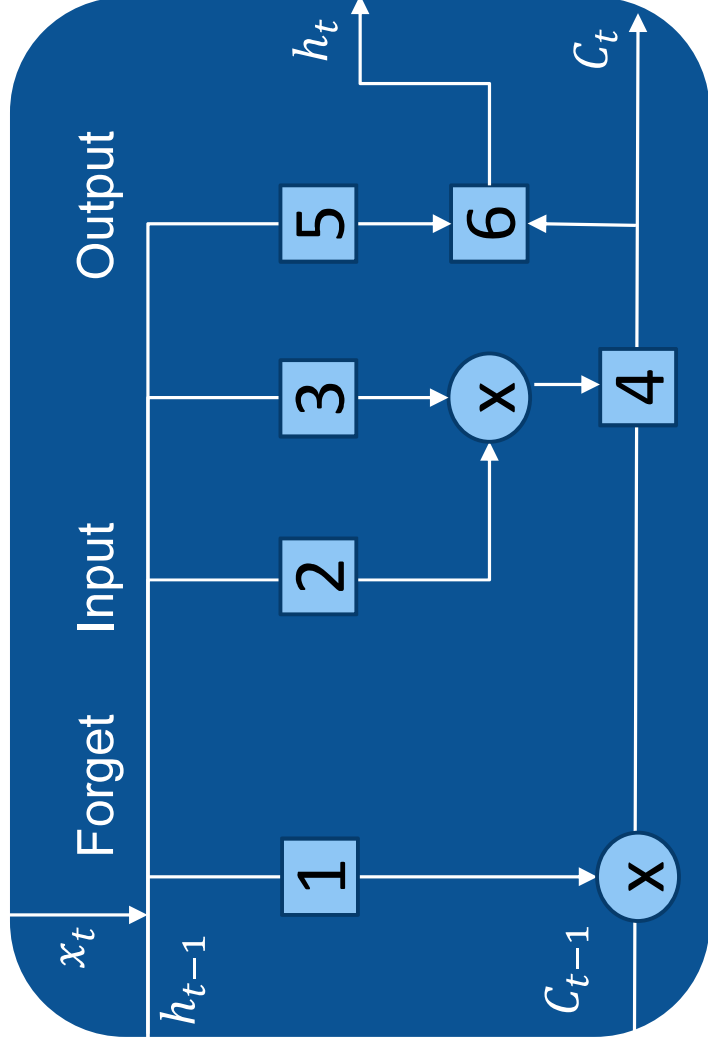
1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Long Short-Term Memory (LSTM)



1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

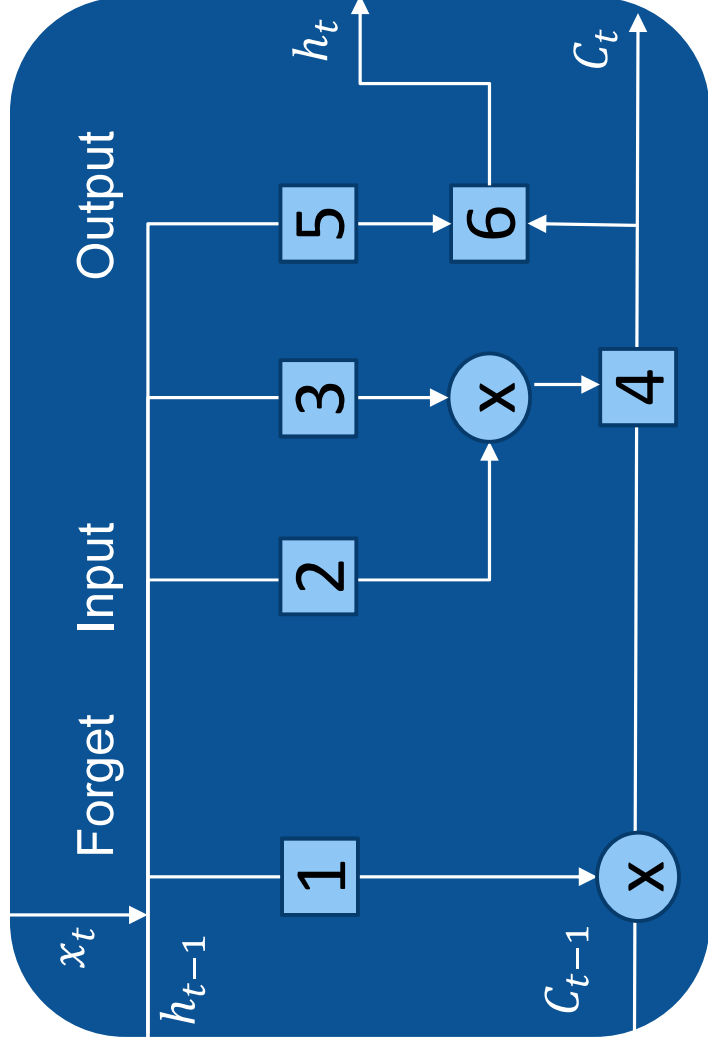
2 $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Long Short-Term Memory (LSTM)



1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

2 $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

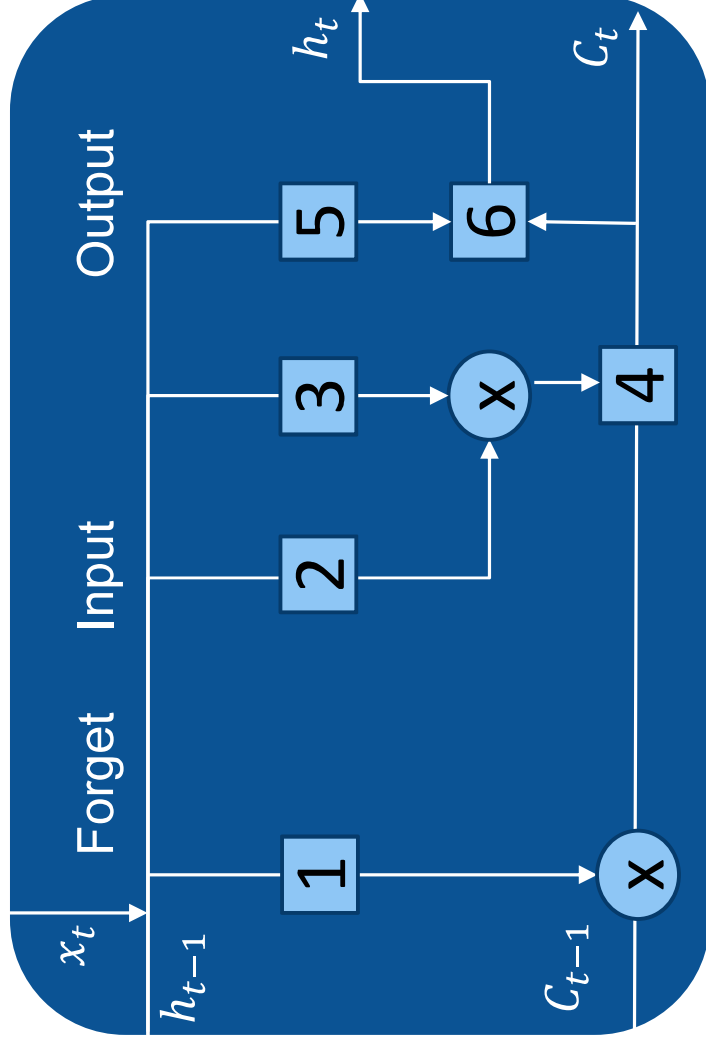
3 $\tilde{C}_t = \tanh(W_C * [h_{t-1}, x_t] + b_C)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Long Short-Term Memory (LSTM)



1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

2 $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

3 $\tilde{C}_t = \tanh(W_C * [h_{t-1}, x_t] + b_C)$

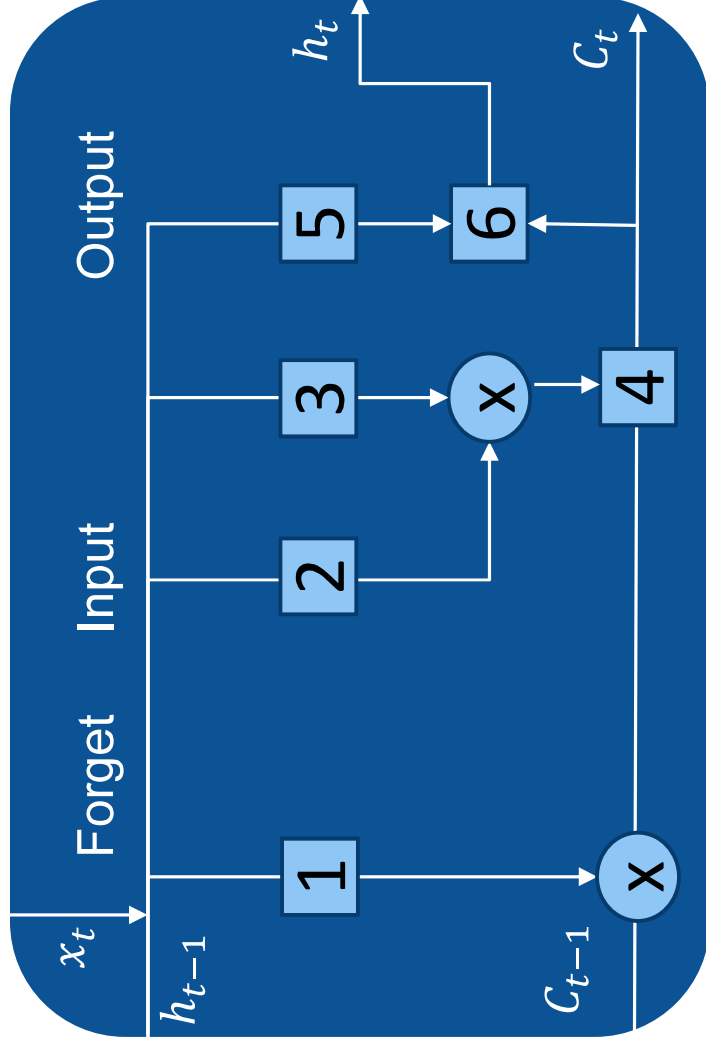
4 $C_t = (f_t * C_{t-1} + i_t * \tilde{C}_t)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Long Short-Term Memory (LSTM)



1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

2 $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

3 $\tilde{C}_t = \tanh(W_C * [h_{t-1}, x_t] + b_C)$

4 $C_t = (f_t * C_{t-1} + i_t * \tilde{C}_t)$

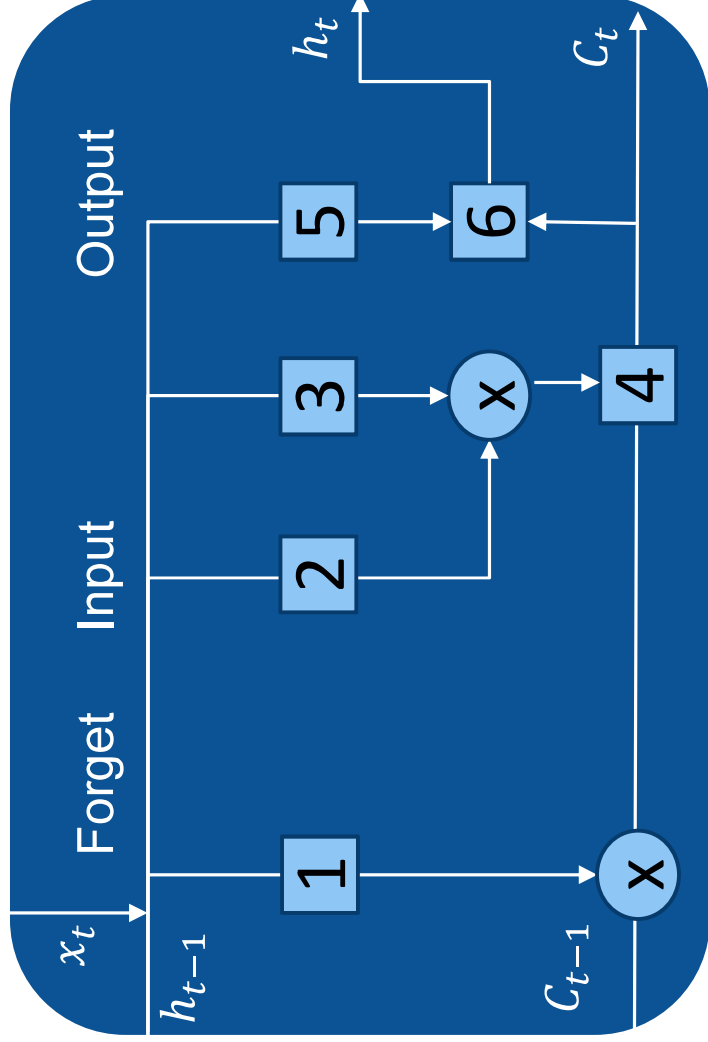
5 $o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Long Short-Term Memory (LSTM)



1 $f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$

2 $i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$

3 $\tilde{C}_t = \tanh(W_C * [h_{t-1}, x_t] + b_C)$

4 $C_t = (f_t * C_{t-1} + i_t * \tilde{C}_t)$

5 $o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$

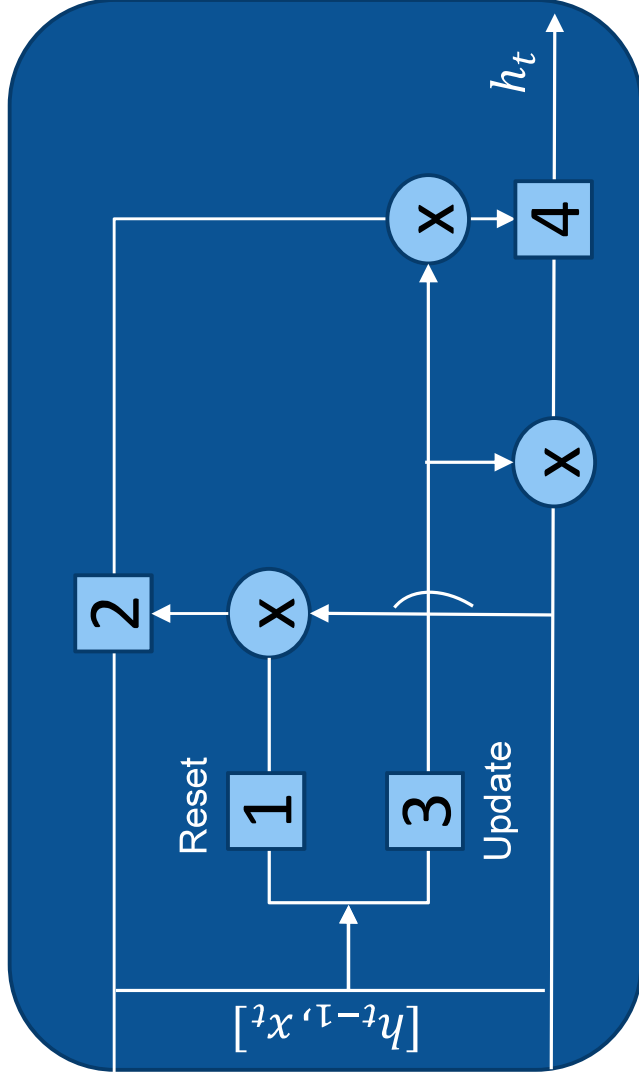
6 $h_t = o_t * \tanh(C_t)$

To solve vanishing/exploding gradient:

- Introduce gates (Input, Output, Forget)



Gated Recurrent Unit



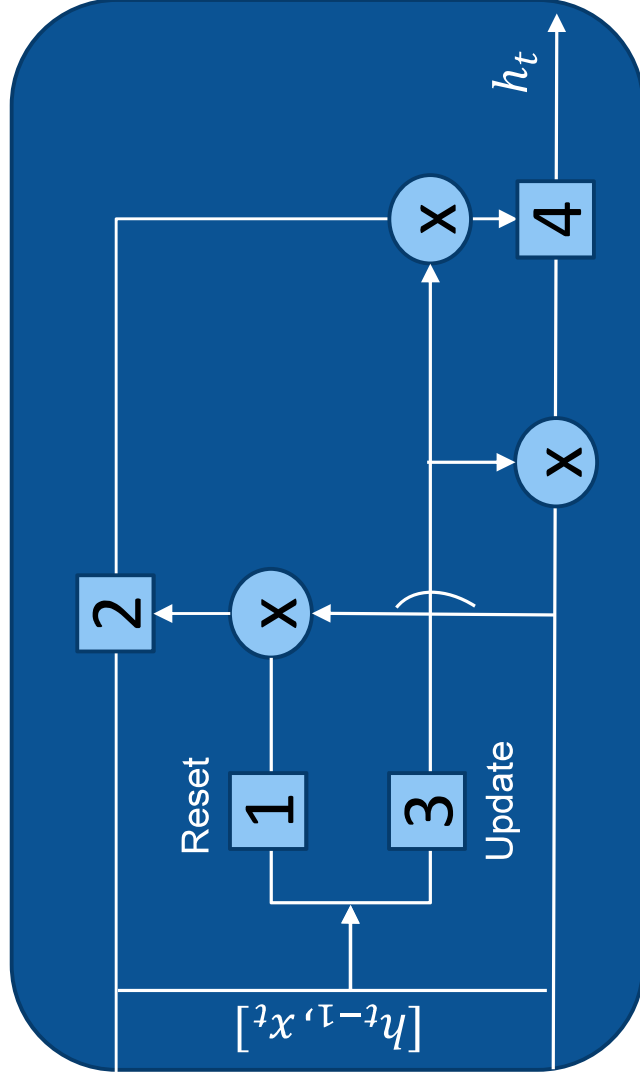
1 $r_t = \sigma(W_r * [h_{t-1}, x_t])$

Key difference with LSTM:

- The forget and input gates are replaced by reset and update gates
- There is no output gate



Gated Recurrent Unit



1

$$r_t = \sigma(W_r * [h_{t-1}, x_t])$$

2

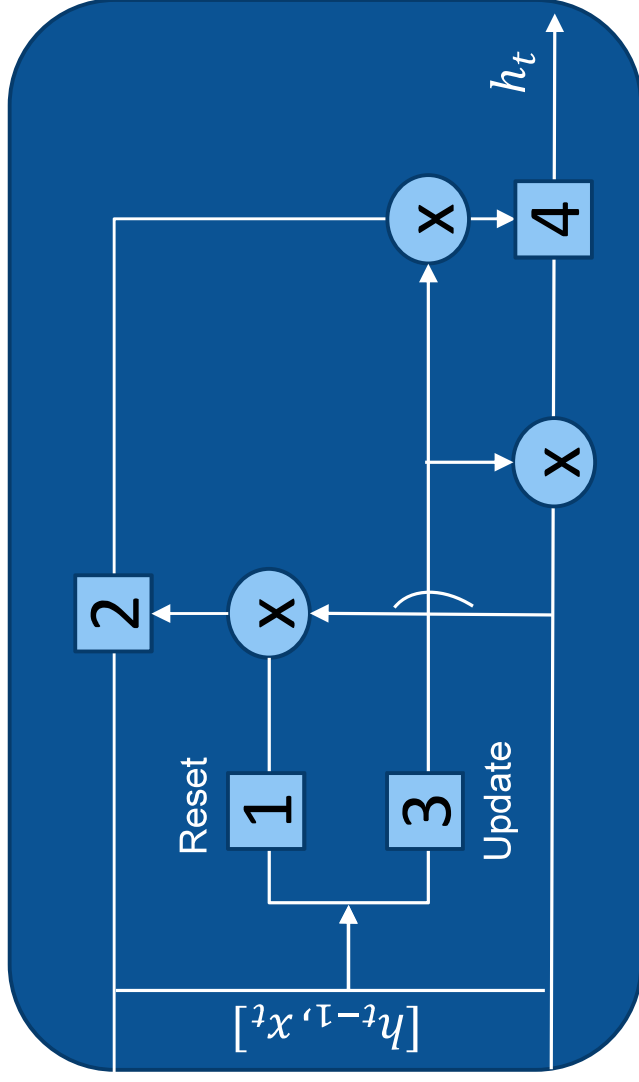
$$\tilde{h}_t = \tanh(W_{\tilde{h}} * [r_t * h_{t-1}, x_t])$$

Key difference with LSTM:

- The forget and input gates are replaced by reset and update gates
- There is no output gate



Gated Recurrent Unit



1

$$r_t = \sigma(W_r * [h_{t-1}, x_t])$$

2

$$\tilde{h}_t = \tanh(W_i * [r_t * h_{t-1}, x_t])$$

3

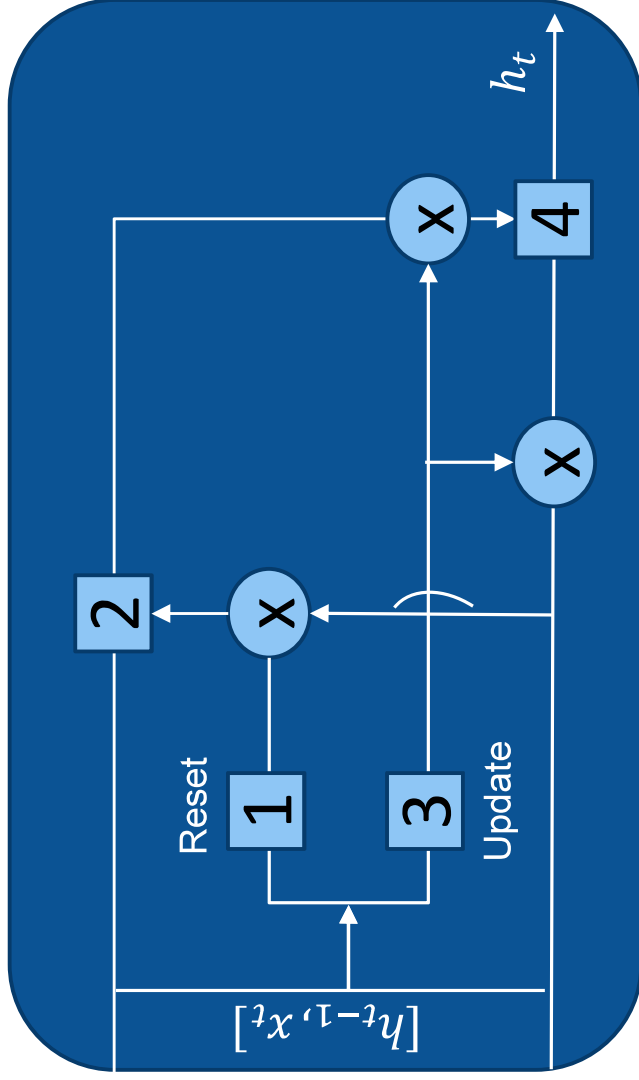
$$z_t = \sigma(W_z * [h_{t-1}, x_t])$$

Key difference with LSTM:

- The forget and input gates are replaced by reset and update gates
- There is no output gate



Gated Recurrent Unit



1

$$r_t = \sigma(W_r * [h_{t-1}, x_t])$$

2

$$\tilde{h}_t = \tanh(W_z * [r_t * h_{t-1}, x_t])$$

3

$$z_t = \sigma(W_z * [h_{t-1}, x_t])$$

4

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Key difference with LSTM:

- The forget and input gates are replaced by reset and update gates
- There is no output gate



Summary

- RNNs handle time-series data by implementing a memory mechanism
- Compared to CNNs are more difficult to train due to the vanishing gradients
- LSTM and GRU are types of RNNs that have been introduced to alleviate the problem of the vanishing/exploding gradients



References

- Ravi et al. Deep Learning for Health Informatics, IEEE Journal of Biomedical and Health Informatics, 21(1), 2017
- Kamath, Deep Learning for NLP Applications, Springer, 2019
- Foster, Generative Deep Learning – Teaching Machines to Paint, Write, Compose and Play, O'Reilly, 2019